

p8106_hw4

Yueyi Xu

2024-04-19

Problem 1

```
# Import data, clean the data, drop na values
college_data <- read.csv("College.csv")

college_data <- college_data %>%
  select(-College)

# Data Partition
set.seed(1)

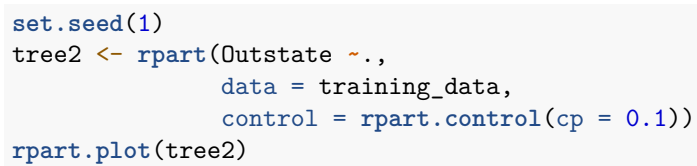
# Split the dataset into training (80%) and test (20%) sets
data_split <- initial_split(college_data, prop = 0.8)

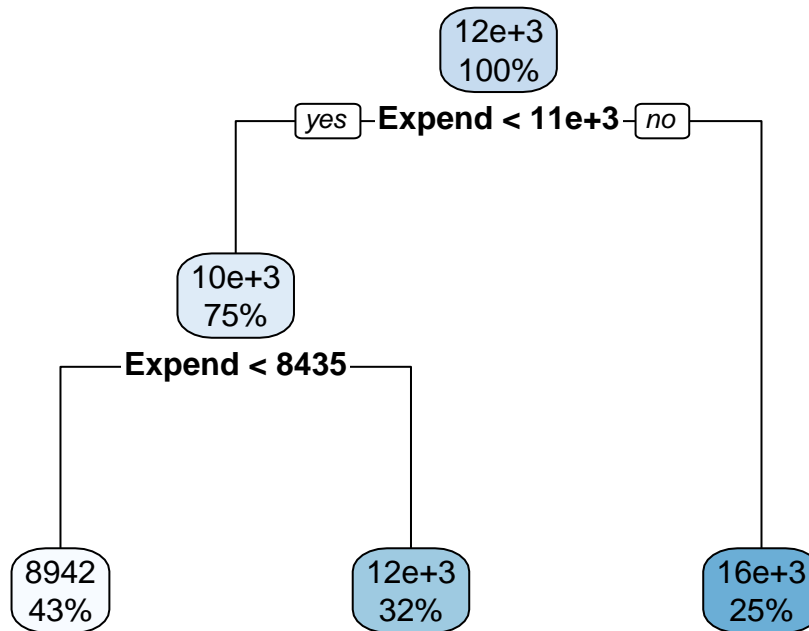
# Extract the training and test data
training_data <- training(data_split)
testing_data <- testing(data_split)
```

(a)

Build a regression tree on the training data to predict the response. Create a plot of the tree.

```
set.seed(1)
tree1 <- rpart(formula = Outstate ~.,
               data = training_data,
               control = rpart.control(cp = 0))
rpart.plot(tree1)
```






```
printcp(tree1)
```

```
##
## Regression tree:
## rpart(formula = Outstate ~ ., data = training_data, control = rpart.control(cp = 0))
##
## Variables actually used in tree construction:
## [1] Accept      Apps      Enroll      Expend      F.Undergrad Grad.Rate
## [7] P.Undergrad perc.alumni Personal  PhD         Room.Board  S.F.Ratio
## [13] Terminal    Top10perc  Top25perc
##
## Root node error: 6.173e+09/452 = 13657034
##
## n= 452
##
##      CP nsplit rel error  xerror   xstd
## 1  0.49687912    0  1.00000 1.00379 0.061129
## 2  0.10633322    1  0.50312 0.51785 0.033714
## 3  0.04791493    2  0.39679 0.44127 0.030355
## 4  0.03504974    3  0.34887 0.38544 0.031017
## 5  0.02506138    4  0.31382 0.36806 0.030846
## 6  0.01516529    5  0.28876 0.36565 0.030603
## 7  0.01094019    6  0.27360 0.36347 0.031752
## 8  0.01077909    7  0.26266 0.36555 0.031353
## 9  0.00959109    8  0.25188 0.35551 0.029394
```

## 10	0.00903094	9	0.24229	0.35649	0.029724
## 11	0.00759360	10	0.23326	0.35001	0.028943
## 12	0.00738206	11	0.22566	0.34463	0.028955
## 13	0.00599129	13	0.21090	0.34313	0.028631
## 14	0.00569487	14	0.20491	0.34252	0.028619
## 15	0.00517513	15	0.19921	0.34177	0.028597
## 16	0.00476813	16	0.19404	0.33579	0.028054
## 17	0.00399063	17	0.18927	0.33458	0.028089
## 18	0.00368136	18	0.18528	0.33666	0.028242
## 19	0.00365828	19	0.18160	0.33573	0.028439
## 20	0.00313676	20	0.17794	0.33417	0.028424
## 21	0.00281892	21	0.17480	0.33610	0.028489
## 22	0.00244723	22	0.17198	0.33976	0.028780
## 23	0.00228629	23	0.16953	0.33979	0.028707
## 24	0.00216557	24	0.16725	0.33917	0.028520
## 25	0.00209368	25	0.16508	0.34306	0.028713
## 26	0.00192948	26	0.16299	0.34334	0.028466
## 27	0.00191531	28	0.15913	0.34274	0.028410
## 28	0.00148413	29	0.15721	0.34481	0.028587
## 29	0.00147516	30	0.15573	0.35026	0.028641
## 30	0.00144328	31	0.15426	0.35172	0.028700
## 31	0.00128466	32	0.15281	0.35314	0.028802
## 32	0.00096241	33	0.15153	0.35236	0.028903
## 33	0.00082504	34	0.15057	0.35435	0.029177
## 34	0.00053977	35	0.14974	0.35088	0.028738
## 35	0.00000000	36	0.14920	0.35046	0.028559

```
cpTable <- tree1$cptable
plotcp(tree1)
```



(b)


Perform random forest on the training data. Report the variable importance and the test error.

```
# Perform random forest
ctrl <- trainControl(method = "cv")

rf.grid <- expand.grid(mtry = 1:16,
                      splitrule = "variance",
                      min.node.size = 1:6)

set.seed(1)
rf.fit <- train(Outstate ~ . ,
               data = training_data,
               method = "ranger",
               tuneGrid = rf.grid,
               trControl = ctrl,
               importance = "permutation")

ggplot(rf.fit, highlight = TRUE) +
  theme_minimal() +
  labs(captain = "Random Forest")
```




```
rf.fit$bestTune %>%
  knitr::kable(caption = "Best tune")
```

Table 1: Best tune

	mtry	splitrule	min.node.size
33	6	variance	3

```
rf2.final.per <- ranger(Outstate ~ . ,
  training_data,
  mtry = rf.fit$bestTune[[1]],
  min.node.size = rf.fit$bestTune[[3]],
  splitrule = "variance",
  importance = "permutation",
  scale.permutation.importance = TRUE)

barplot(sort(ranger::importance(rf2.final.per),
  decreasing = FALSE),
  las = 2,
  horiz = TRUE,
  cex.names = 0.7,
  col = colorRampPalette(colors = c("cyan", "blue"))(19))
```



```
# Variable Importance
df <- as.data.frame(ranger::importance(rf2.final.per))
colnames(df) <- "Importance"
df %>% arrange(desc(Importance))
```

```
##           Importance
## Expend      40.178942
## Room.Board  29.081572
## Apps        20.346272
## Terminal    16.534106
## Accept      14.196601
## perc.alumni 13.800836
## Grad.Rate   12.774939
## Top10perc   12.368317
## Enroll      12.290758
## PhD         11.475855
## S.F.Ratio   11.183291
## Top25perc   11.077537
## P.Undergrad 10.798701
## F.Undergrad 10.537615
## Personal     9.616744
## Books       2.704511
```

```
# Test Error
pred.rf <- predict(rf.fit, newdata = testing_data)
RMSE.rf <- RMSE(pred.rf, testing_data$Outstate)
```

The variable “expend” is the most significant variable with 40.178942 value for predicting out-of-state tuition. The test error of RMSE for the random forest model is 1741.258.


(c)

Perform boosting on the training data. Report the variable importance and the test error.

```
# Boosting
gbm.grid <- expand.grid(n.trees = c(5000,10000,20000,30000,40000,50000),
  interaction.depth = 1:3,
  shrinkage = c(0.001,0.005),
  n.minobsinnode = c(1))

set.seed(1)
gbm.fit <- train(Outstate ~ . ,
  data = training_data,
  method = "gbm",
  tuneGrid = gbm.grid,
  trControl = ctrl,
  verbose = FALSE)

ggplot(gbm.fit, highlight = TRUE) +
  theme_minimal() +
  labs(caption = "Boosting")
```




```
gbm.fit$bestTune %>%
  knitr::kable(caption = "Best tune for boostin")
```

Table 2: Best tune for boostin

	n.trees	interaction.depth	shrinkage	n.minobsinnode
14	10000	3	0.001	1

```
# Variable Importance
summary(gbm.fit$finalModel, plot = FALSE)
```

```
##           var  rel.inf
## Expend      Expend 53.244888
## Room.Board  Room.Board 10.230882
## Terminal    Terminal  5.940624
## Grad.Rate    Grad.Rate  4.278428
## perc.alumni  perc.alumni 3.502031
## Apps         Apps      3.387694
## F.Undergrad F.Undergrad 2.410000
## PhD          PhD       2.313853
## Personal     Personal  2.292650
## P.Undergrad  P.Undergrad 2.164553
## Accept       Accept    2.145142
## S.F.Ratio    S.F.Ratio  2.019490
## Top25perc    Top25perc  1.921880
## Top10perc    Top10perc  1.694928
## Books        Books     1.251389
## Enroll       Enroll    1.201568
```

```
# Test Error
gbm.pred <- predict(gbm.fit, newdata = testing_data)
RMSE.boosting <- RMSE(gbm.pred, testing_data$Outstate)
```

The variable “expend” is the most significant variable with 53.244888 value for predicting out-of-state tuition. The test error of RMSE for the boosting model is 1649.905.

Problem 2

```
auto_data <- read.csv("auto.csv")

auto_data <- auto_data %>%
  mutate(mpg_cat = as.factor(mpg_cat))

set.seed(1)
data_split_auto <- initial_split(auto_data, prop = 0.7)
# Extract the training and testing data
training_data_auto <- training(data_split_auto)
testing_data_auto <- testing(data_split_auto)
```

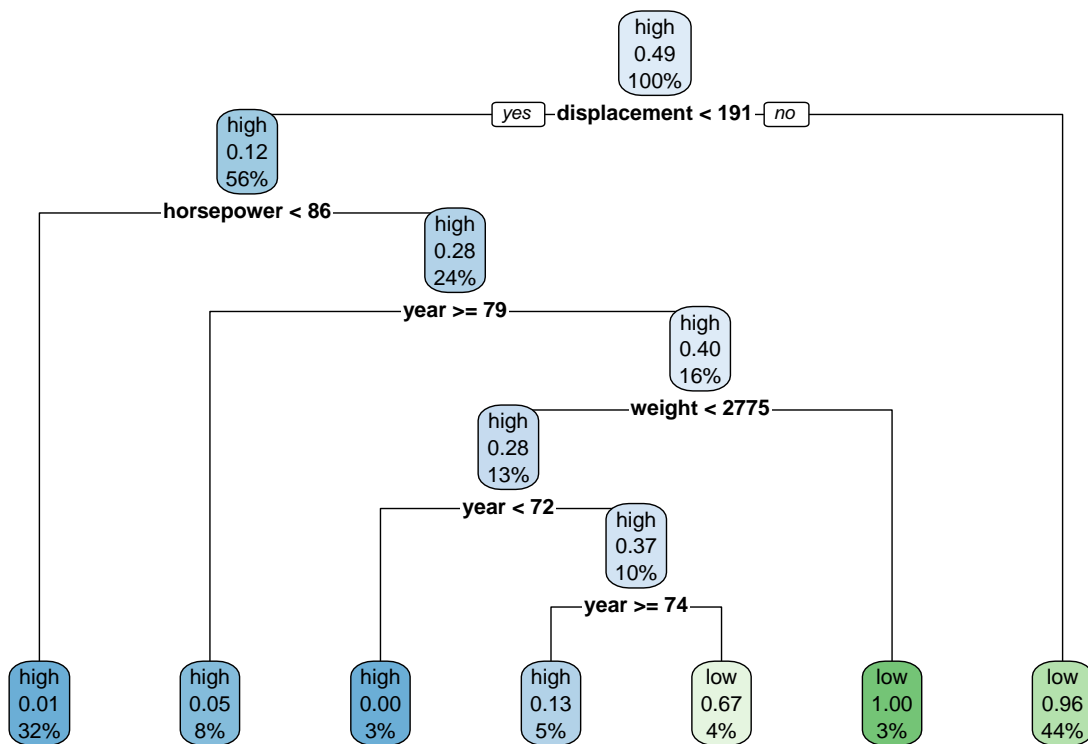
(a)

Build a classification tree using the training data, with mpg cat as the response. Which tree size corresponds to the lowest cross-validation error? Is this the same as the tree size obtained using the 1 SE rule?

```
# Classification tree
ctrl_2 <- trainControl(method = "cv",
  summaryFunction = twoClassSummary,
  classProbs = TRUE)

set.seed(1)
rpart.fit <- train(mpg_cat ~ . ,
  training_data_auto,
  method = "rpart",
  tuneGrid = data.frame(cp = exp(seq(-8, -3, len = 100))),
  trControl = ctrl_2,
  metric = "ROC")

best.ct <- prune(rpart.fit$finalModel, cp = rpart.fit$bestTune$cp)
rpart.plot(best.ct)
```



```
size1 <- nrow(best.ct$frame)
rpart.fit$bestTune %>%
  knitr::kable(caption = "Best tune for classification tree")
```

Table 3: Best tune for classification tree

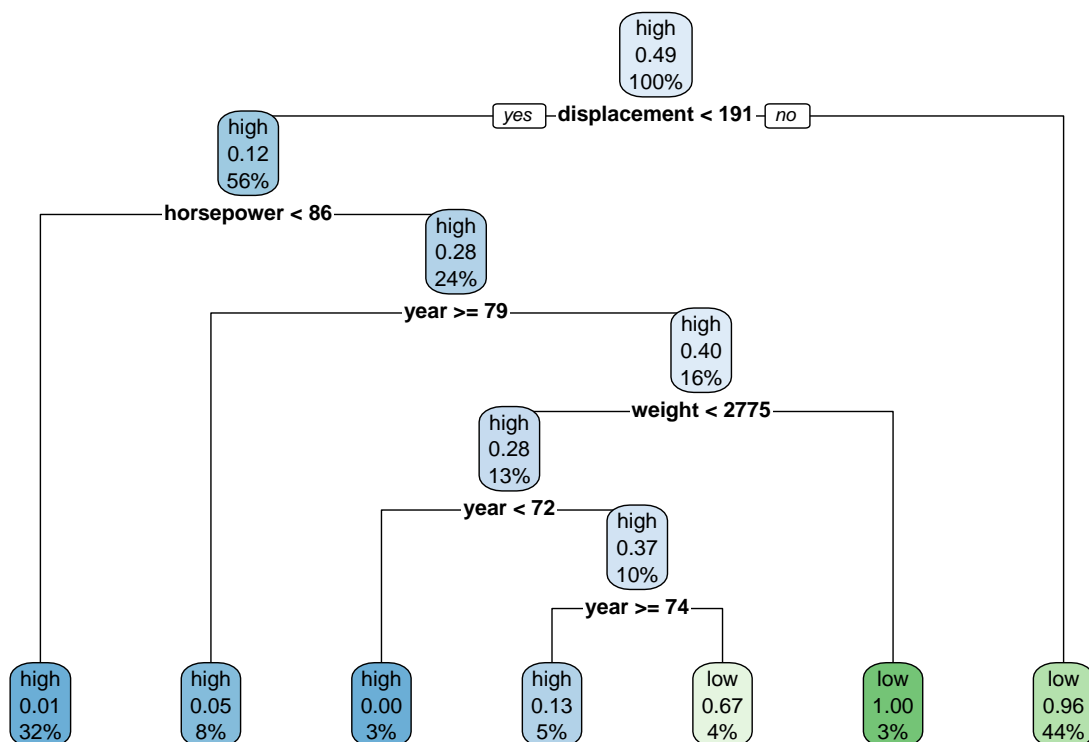
	cp
50	0.0039849

```
# 1SE
ctrl_1se <- trainControl(method = "cv",
  summaryFunction = twoClassSummary,
  selectionFunction = "oneSE",
  classProbs = TRUE)

set.seed(1)

rpart.fit_1se <- train(mpg_cat ~ . ,
  training_data_auto,
  method = "rpart",
  tuneGrid = data.frame(cp = exp(seq(-8,-3, len = 100))),
  trControl = ctrl_1se,
  metric = "ROC")

oneSE.ct <- prune(rpart.fit_1se$finalModel, cp = rpart.fit_1se$bestTune$cp)
rpart.plot(oneSE.ct)
```



```
size2 <- nrow(oneSE.ct$frame)
rpart.fit_1se$bestTune %>%
  knitr::kable(caption = "Best tune for 1SE classification tree")
```

Table 4: Best tune for 1SE classification tree

	cp
64	0.0080815

```
# Displaying best cp values
```

```
cat("Best cp without 1SE rule:", rpart.fit$bestTune$cp, "\n")
```

```
## Best cp without 1SE rule: 0.003984862
```

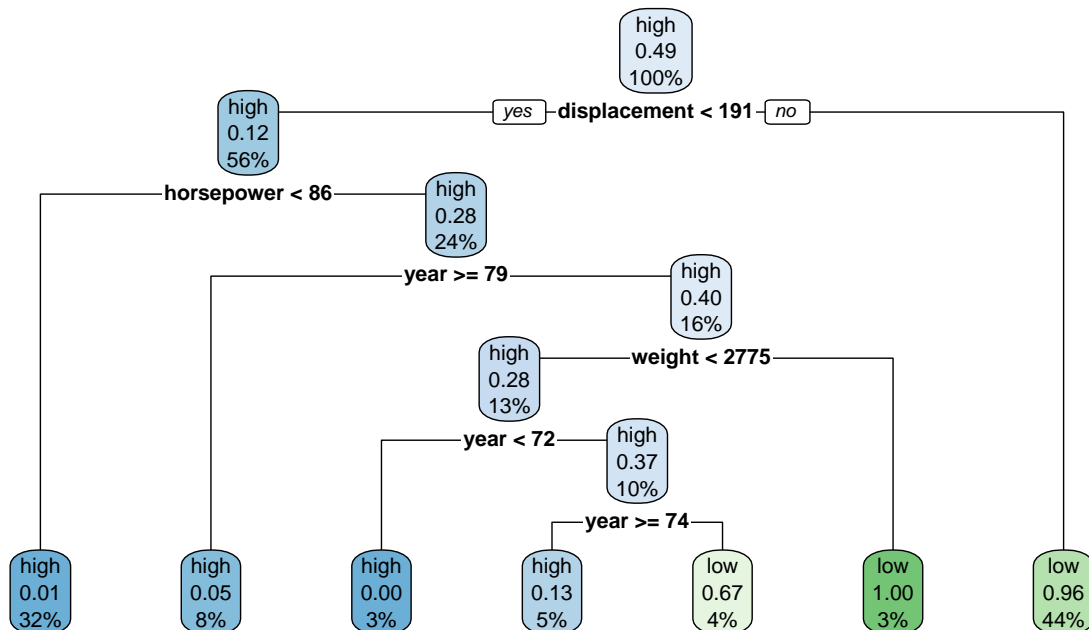
```
cat("Best cp with 1SE rule:", rpart.fit_1se$bestTune$cp, "\n")
```

```
## Best cp with 1SE rule: 0.008081467
```

```
# Compare trees
```

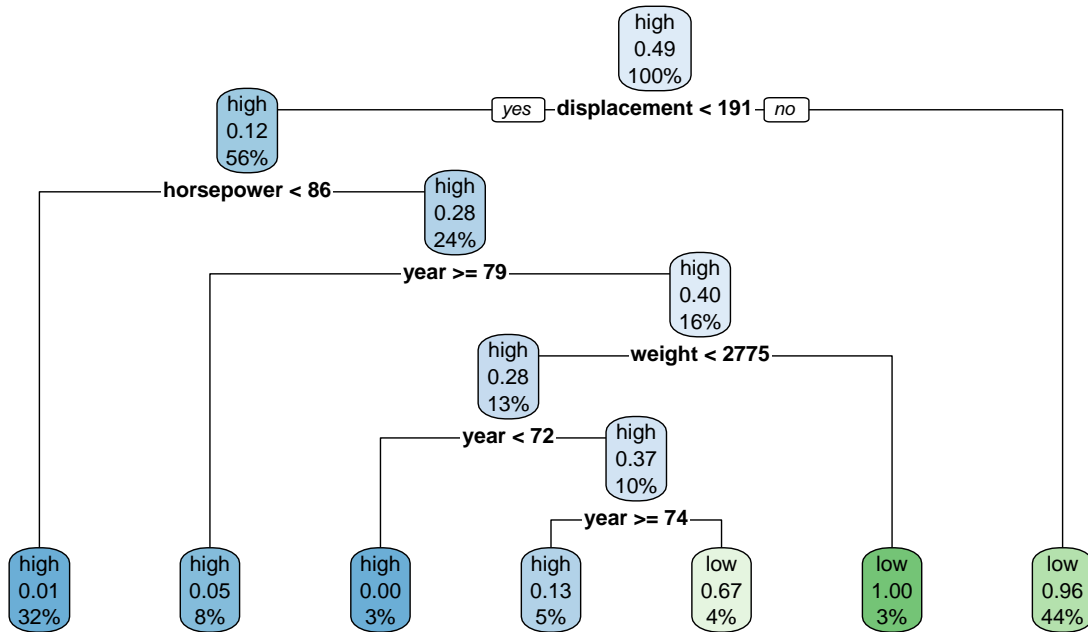
```
rpart.plot(rpart.fit$finalModel, main = "Best Model without 1SE rule")
```

Best Model without 1SE rule



```
rpart.plot(rpart.fit_1se$finalModel, main = "Best Model with 1SE rule")
```

Best Model with 1SE rule



Both without the 1SE rule and with the 1SE rule, the tree size is 13. However, the best model without the 1SE rule has the complexity parameter of 0.003984862 while the model with 1SE has the complexity parameter of 0.008081467.

(b)

Perform boosting on the training data and report the variable importance. Report the test data performance.

```

# Boosting
ctrl_3 <- trainControl(method = "cv",
  summaryFunction = twoClassSummary,
  classProbs = TRUE,
  selectionFunction = "best")

gbm.grid.auto <- expand.grid(n.trees = c(5000,10000,20000,30000,40000,50000),
  interaction.depth = 1:3,
  shrinkage = c(0.001,0.005),
  n.minobsinnode = c(1))

set.seed(1)
gbm.fit.auto <- train(mpg_cat ~ . ,
  data = training_data_auto,
  method = "gbm",
  tuneGrid = gbm.grid,
  trControl = ctrl_3,
  verbose = FALSE)
  
```

```
## Warning in train.default(x, y, weights = w, ...): The metric "Accuracy" was not
## in the result set. ROC will be used instead.
```

```
# Variable Importance
summary(gbm.fit.auto$finalModel, plot = FALSE)
```

```
##              var      rel.inf
## displacement displacement 39.21302364
## weight          weight 23.20004110
## cylinders        cylinders 18.90377862
## horsepower       horsepower 12.90864349
## year            year  5.19847041
## acceleration     acceleration 0.54678952
## origin           origin 0.02925322
```

```
# Test Error
gbm.pred.auto <- predict(gbm.fit.auto, newdata = testing_data_auto, type = 'prob')[,1]
pROC::roc(testing_data_auto$mpg_cat, gbm.pred.auto)
```

```
## Setting levels: control = high, case = low
```

```
## Setting direction: controls > cases
```

```
##
```

```
## Call:
```

```
## roc.default(response = testing_data_auto$mpg_cat, predictor = gbm.pred.auto)
```

```
##
```

```
## Data: gbm.pred.auto in 57 controls (testing_data_auto$mpg_cat high) > 61 cases (testing_data_auto$mpg_cat low)
```

```
## Area under the curve: 0.9781
```

The variable “displacement” is the most significant variable with 39.21302364 value for predicting out-of-state tuition. The AUC value is 0.9781.